

ProcessingV2

April 28, 2024

```
[26]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import scipy
import sklearn
import os
```

```
[27]: def find_csv_files(folder_path, suffix=".csv"):
    """
    Find all CSV files in the given folder path.

    Args:
    - folder_path (str): Path to the folder where CSV files are located.
    - suffix (str): Suffix to filter files (default is '.csv').

    Returns:
    - list: List of CSV files with the given suffix.
    """
    csv_files = []
    for root, dirs, files in os.walk(folder_path):
        for file in files:
            if file.endswith(suffix):
                csv_files.append(os.path.join(root, file))
    return csv_files
```

```
[28]: def plot_top_n_categorical(data, column, n=None, figsize=(15, 5)):
    """
    Plot the distribution of the top N categories of a categorical variable.

    Args:
    - data (pandas DataFrame): DataFrame containing the categorical variable.
    - column (str): Name of the categorical variable column.
    - n (int or None): Number of top categories to plot. If None, plot all
    ↪ categories (default is None).
    - figsize (tuple): Width and height of the figure in inches (default is
    ↪ (10, 6)).
```

```

Returns:
- None (displays the plot).
"""
if n is None:
    categories = data[column].value_counts().index
else:
    categories = data[column].value_counts().nlargest(n).index

data_filtered = data[data[column].isin(categories)]

plt.figure(figsize=figsize)
sns.countplot(data=data_filtered, x=column, order=categories)
plt.title(f'Top {len(categories)} Categories of {column}')
plt.xlabel(column)
plt.ylabel('Count')
plt.xticks(rotation=90)
plt.show()

```

```
[29]: data = pd.read_csv("./INITIAL_PROCESSED_DATA.csv")
```

```
[30]: data = data.sort_values(by=['Date (MM/DD/YYYY)', 'Scheduled Arrival Time'])
```

```
[31]: to_drop = [
    'Arrival Delay (Minutes)',
    'Flight Number',
    'Tail Number',
]
```

```
[32]: data = data.drop(columns=to_drop)
```

```
[33]: weather_data_path = './imputed_weather/{ }_weather_data.csv'
```

```
[34]: data.head()
```

```
[34]:
```

	Carrier Code	Date (MM/DD/YYYY)	Origin Airport	Scheduled Arrival Time	\
0	B6	2010-01-01	JFK	00:01	
1	B6	2010-01-01	JFK	08:55	
2	MQ	2010-01-01	ORD	11:20	
3	9E	2010-01-01	DTW	11:44	
4	B6	2010-01-01	JFK	11:52	

	Scheduled Elapsed Time (Minutes)	FLIGHT_STATUS	month	day	season	WeekDay
0	76	LATE	1	1	winter	Friday
1	75	LATE	1	1	winter	Friday
2	100	ONTIME	1	1	winter	Friday
3	84	LATE	1	1	winter	Friday

4	71	LATE	1	1	winter	Friday
---	----	------	---	---	--------	--------

```
[35]: data['Scheduled Arrival Time'] = data['Scheduled Arrival Time'].replace('24:
↳00', '23:59')
```

```
[36]: data['UNIX_DATE'] = pd.to_datetime(data['Date (MM/DD/YYYY)'] + ' ' +
↳data['Scheduled Arrival Time'])
```

```
[37]: data['UNIX_TIMESTAMP'] = data['UNIX_DATE'].apply(lambda x: int(x.timestamp()))
```

```
[38]: data.head()
```

```
[38]:  Carrier Code Date (MM/DD/YYYY) Origin Airport Scheduled Arrival Time \
0      B6      2010-01-01      JFK      00:01
1      B6      2010-01-01      JFK      08:55
2      MQ      2010-01-01      ORD      11:20
3      9E      2010-01-01      DTW      11:44
4      B6      2010-01-01      JFK      11:52

      Scheduled Elapsed Time (Minutes) FLIGHT_STATUS month day season WeekDay \
0      76      LATE      1      1      winter      Friday
1      75      LATE      1      1      winter      Friday
2     100     ONTIME      1      1      winter      Friday
3      84      LATE      1      1      winter      Friday
4      71      LATE      1      1      winter      Friday

      UNIX_DATE UNIX_TIMESTAMP
0 2010-01-01 00:01:00      1262304060
1 2010-01-01 08:55:00      1262336100
2 2010-01-01 11:20:00      1262344800
3 2010-01-01 11:44:00      1262346240
4 2010-01-01 11:52:00      1262346720
```

1 Weather

```
[39]: def mode_imputer(data):
      """
      Perform mode imputation on a DataFrame.

      Parameters:
          data (DataFrame): Input DataFrame with missing values.

      Returns:
          DataFrame: DataFrame with missing values replaced by mode.
      """
      # Fill missing values with mode
```

```
data_imputed = data.fillna(data.mode().iloc[0])

return data_imputed
```

```
[40]: all_weather_data = find_csv_files('./imputed_weather/')
```

```
[ ]:
```

```
[41]: # for d in all_weather_data:
#     dd = mode_imputer(pd.read_csv(d))
#     print(f"===== {d} =====")
#     dd = dd.fillna(1013)
#     print(dd.isna().sum())
#     print("=====")
```

```
[42]: weather1 = pd.read_csv(all_weather_data[0], parse_dates=['date'])
```

```
[43]: weather1['UNIX_TIMESTAMP'] = weather1['date'].apply(lambda x: int(x.
↳ timestamp()))
```

```
[44]: weather1.head()
```

```
[44]: station          date  latitude  longitude  elevation  wind_direction \
0      MSP 2009-01-01 00:00:00  44.8831   -93.2289      265.8        140.0
1      MSP 2009-01-01 00:53:00  44.8831   -93.2289      265.8        140.0
2      MSP 2009-01-01 01:53:00  44.8831   -93.2289      265.8        150.0
3      MSP 2009-01-01 02:29:00  44.8831   -93.2289      265.8        140.0
4      MSP 2009-01-01 02:53:00  44.8831   -93.2289      265.8        160.0
```

```
wind_type  wind_speed  ceiling_height  ceiling_det_code  celing_CAVOK \
0          N         36.0      3309.428571              9          N
1          N         36.0      3658.000000              M          N
2          V         31.0      3048.000000              M          N
3          N         62.0      3353.000000              M          N
4          N         46.0      3048.000000              M          N
```

```
visibility_dist  visibility_variability  air_temperature \
0          16000.0                  N          -156.0
1          16093.0                  N          -156.0
2          16093.0                  N          -133.0
3          16093.0                  N          -130.0
4          16093.0                  N          -128.0
```

```
dew_point_temperature  sea_level_pressure  UNIX_TIMESTAMP
0          -211.0          10274.0      1230768000
1          -206.0          10255.0      1230771180
2          -194.0          10237.0      1230774780
```

3	-190.0	10204.8	1230776940
4	-189.0	10217.0	1230778380

```
[45]: def binary_search_nearest_rows(data, target_timestamp, num_neighbors=3):
    """
    Perform binary search to find the nearest rows in a DataFrame for a given
    ↪ timestamp.

    Parameters:
        data (DataFrame): Input DataFrame with timestamps.
        target_timestamp (int): Target timestamp for which the nearest rows
    ↪ need to be found.
        num_neighbors (int): Number of nearest rows to find on each side of the
    ↪ target timestamp.

    Returns:
        DataFrame: DataFrame containing the rows corresponding to the nearest
    ↪ timestamps.
    """
    # Sort DataFrame by the timestamp column
    data_sorted = data.sort_values(by='UNIX_TIMESTAMP')

    # Convert the sorted timestamps column to a list
    timestamps = data_sorted['UNIX_TIMESTAMP'].tolist()

    # Binary search to find the nearest timestamp
    low = 0
    high = len(timestamps) - 1

    while low <= high:
        mid = (low + high) // 2
        mid_timestamp = timestamps[mid]

        if mid_timestamp == target_timestamp:
            nearest_indices = [mid]
            break

        elif mid_timestamp < target_timestamp:
            low = mid + 1

        else:
            high = mid - 1

    else:
        # Find the nearest timestamp
        nearest_indices = []
        if high < 0:
```

```

        nearest_indices.append(low)
    elif low >= len(timestamps):
        nearest_indices.append(high)
    else:
        if abs(timestamps[low] - target_timestamp) < abs(timestamps[high] -
↪target_timestamp):
            nearest_indices.append(low)
        else:
            nearest_indices.append(high)

    # Ensure we have enough neighbors
    while len(nearest_indices) < num_neighbors:
        if nearest_indices[0] - 1 >= 0:
            nearest_indices.insert(0, nearest_indices[0] - 1)
        else:
            nearest_indices.append(nearest_indices[-1] + 1)

    # Extract the rows corresponding to the nearest timestamps
    nearest_rows = data_sorted.iloc[nearest_indices]

    return nearest_rows

```

```

[46]: binary_search_nearest_rows(weather1, 1262346720)[['latitude', 'longitude',
↪'elevation',
        'wind_direction', 'wind_type', 'wind_speed', 'ceiling_height',
        'ceiling_det_code', 'celing_CAVOK', 'visibility_dist',
        'visibility_variability', 'air_temperature', 'dew_point_temperature',
        'sea_level_pressure', 'UNIX_TIMESTAMP']]

```

```

[46]:
    latitude  longitude  elevation  wind_direction  wind_type  wind_speed  \
12994    44.8831   -93.2289     265.8           300.0         N         31.0
12995    44.8831   -93.2289     265.8           320.0         N         36.0
12996    44.8831   -93.2289     265.8           310.0         N         21.0

    ceiling_height  ceiling_det_code  celing_CAVOK  visibility_dist  \
12994           366.0                M            N           16093.0
12995           610.0                M            N           16093.0
12996           701.0                M            N           16093.0

    visibility_variability  air_temperature  dew_point_temperature  \
12994                    N             -170.0             -210.0
12995                    N             -170.0             -210.0
12996                    N             -170.0             -210.0

    sea_level_pressure  UNIX_TIMESTAMP
12994      10302.932552      1262345460
12995      10306.167846      1262345880

```

12996

10309.400485

1262346660

2 FINAL DATA PREP

```
[48]: import tqdm
all_weather = {}
all_rows = []
SYR_DF = pd.read_csv(weather_data_path.format('SYR'), parse_dates=['date'])
SYR_DF['UNIX_TIMESTAMP'] = SYR_DF['date'].apply(lambda x: int(x.timestamp()))
for index, row in tqdm.tqdm_notebook(data.iterrows(), total=len(data)):
    station = row['Origin Airport']
    if station not in all_weather:
        all_weather[station] = pd.read_csv(weather_data_path.format(station),
        ↪ parse_dates=['date'])
        all_weather[station]['UNIX_TIMESTAMP'] = all_weather[station]['date'].
        ↪ apply(lambda x: int(x.timestamp()))
        weatherDF = all_weather[station]
        row = row.to_dict()
        row_unix_ts = row['UNIX_TIMESTAMP']

        closest_rows = binary_search_nearest_rows(weatherDF, row_unix_ts,
        ↪ num_neighbors=5)[['latitude', 'longitude', 'elevation',
            'wind_direction', 'wind_type', 'wind_speed', 'ceiling_height',
            'ceiling_det_code', 'celing_CAVOK', 'visibility_dist',
            'visibility_variability', 'air_temperature', 'dew_point_temperature',
            'sea_level_pressure']]

        for nc in ['latitude', 'longitude', 'elevation',
            'wind_direction', 'wind_type', 'wind_speed', 'ceiling_height',
            'ceiling_det_code', 'celing_CAVOK', 'visibility_dist',
            'visibility_variability', 'air_temperature', 'dew_point_temperature',
            'sea_level_pressure']:
            if nc in ['wind_type', 'ceiling_det_code', 'celing_CAVOK',
            ↪ 'visibility_variability']:
                row[nc] = closest_rows[nc].mode().iloc[0]
            else:
                row[nc] = closest_rows[nc].mean()

        closest_rows = binary_search_nearest_rows(SYR_DF, row_unix_ts,
        ↪ num_neighbors=5)[['latitude', 'longitude', 'elevation',
            'wind_direction', 'wind_type', 'wind_speed', 'ceiling_height',
            'ceiling_det_code', 'celing_CAVOK', 'visibility_dist',
            'visibility_variability', 'air_temperature', 'dew_point_temperature',
            'sea_level_pressure']]

        for nc in ['latitude', 'longitude', 'elevation',
```

```

        'wind_direction', 'wind_type', 'wind_speed', 'ceiling_height',
        'ceiling_det_code', 'celing_CAVOK', 'visibility_dist',
        'visibility_variability', 'air_temperature', 'dew_point_temperature',
        'sea_level_pressure']:
        if nc in ['wind_type', 'ceiling_det_code', 'celing_CAVOK',
        ↪ 'visibility_variability']:
            row['SYR_'+nc] = closest_rows[nc].mode().iloc[0]
        else:
            row['SYR_'+nc] = closest_rows[nc].mean()
        all_rows.append(row)
    new_data_df = pd.DataFrame(all_rows, index=data.index)

```

/tmp/ipykernel_32272/3525575574.py:6: TqdmDeprecationWarning: This function will be removed in tqdm==5.0.0

Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`

```
for index, row in tqdm.tqdm_notebook(data.iterrows(), total=len(data)):
```

```
0%|          | 0/113671 [00:00<?, ?it/s]
```

[49]: new_data_df.head()

```

[49]:  Carrier Code Date (MM/DD/YYYY) Origin Airport Scheduled Arrival Time \
0      B6      2010-01-01      JFK      00:01
1      B6      2010-01-01      JFK      08:55
2      MQ      2010-01-01      ORD      11:20
3      9E      2010-01-01      DTW      11:44
4      B6      2010-01-01      JFK      11:52

      Scheduled Elapsed Time (Minutes) FLIGHT_STATUS  month  day  season WeekDay \
0      76      LATE      1    1  winter  Friday
1      75      LATE      1    1  winter  Friday
2     100     ONTIME      1    1  winter  Friday
3      84      LATE      1    1  winter  Friday
4      71      LATE      1    1  winter  Friday

      ... SYR_wind_type  SYR_wind_speed  SYR_ceiling_height \
0  ...      N      17.4      741.370199
1  ...      C       0.0      548.800000
2  ...      C       0.0      841.400000
3  ...      C       0.0     1006.000000
4  ...      C       0.0     1006.000000

      SYR_ceiling_det_code  SYR_celing_CAVOK  SYR_visibility_dist \
0      M      N      6110.0
1      M      N      2414.0
2      M      N      3138.4
3      M      N      3138.4
4      M      N      3138.4

```


	SYR_visibility_variability	SYR_air_temperature	SYR_dew_point_temperature \
0	N	-4.8	-24.0
1	N	-17.0	-28.8
2	N	-19.4	-29.6
3	N	-19.8	-29.2
4	N	-19.8	-29.2

	SYR_sea_level_pressure
0	10154.572817
1	10132.352069
2	10129.796558
3	10129.482076
4	10129.482076

[5 rows x 40 columns]

```
[50]: new_data_df.isna().sum()
```

```
[50]: Carrier Code          0
Date (MM/DD/YYYY)          0
Origin Airport              0
Scheduled Arrival Time      0
Scheduled Elapsed Time (Minutes) 0
FLIGHT_STATUS               0
month                       0
day                         0
season                     0
WeekDay                    0
UNIX_DATE                  0
UNIX_TIMESTAMP             0
latitude                   0
longitude                  0
elevation                  0
wind_direction             0
wind_type                  0
wind_speed                 0
ceiling_height             0
ceiling_det_code           0
celing_CAVOK               0
visibility_dist            0
visibility_variability      0
air_temperature            0
dew_point_temperature       0
sea_level_pressure         728
SYR_latitude               0
SYR_longitude              0
```

SYR_elevation	0
SYR_wind_direction	0
SYR_wind_type	0
SYR_wind_speed	0
SYR_ceiling_height	0
SYR_ceiling_det_code	0
SYR_celing_CAVOK	0
SYR_visibility_dist	0
SYR_visibility_variability	0
SYR_air_temperature	0
SYR_dew_point_temperature	0
SYR_sea_level_pressure	0
dtype: int64	

```
[51]: new_data_df['sea_level_pressure'] = new_data_df['sea_level_pressure'].
      ↪ fillna(new_data_df['sea_level_pressure'].mean())
```

```
[52]: new_data_df.isna().sum()
```

[52]: Carrier Code	0
Date (MM/DD/YYYY)	0
Origin Airport	0
Scheduled Arrival Time	0
Scheduled Elapsed Time (Minutes)	0
FLIGHT_STATUS	0
month	0
day	0
season	0
WeekDay	0
UNIX_DATE	0
UNIX_TIMESTAMP	0
latitude	0
longitude	0
elevation	0
wind_direction	0
wind_type	0
wind_speed	0
ceiling_height	0
ceiling_det_code	0
celing_CAVOK	0
visibility_dist	0
visibility_variability	0
air_temperature	0
dew_point_temperature	0
sea_level_pressure	0
SYR_latitude	0
SYR_longitude	0

```

SYR_elevation                0
SYR_wind_direction           0
SYR_wind_type                 0
SYR_wind_speed                0
SYR_ceiling_height            0
SYR_ceiling_det_code          0
SYR_celing_CAVOK              0
SYR_visibility_dist           0
SYR_visibility_variability     0
SYR_air_temperature           0
SYR_dew_point_temperature     0
SYR_sea_level_pressure        0
dtype: int64

```

```
[60]: # new_data_df.to_csv("SYR_ORIGIN_WEATHER_IMPUTED.csv", index=False)
```

```
[61]: def create_data_2nd_model_data_ext(df, num_prev=3):
    extended_data = []
    for index, row in df.iterrows():
        for i in range(num_prev):
            if index - i - 1 >= 0:
                extended_row = row.copy() # Create a copy of the current row
                extended_row['PREV_STAT'] = df.iloc[index - i - 1]
                extended_data.append(extended_row)
            else:
                extended_row = row.copy() # Create a copy of the current row
                extended_row['PREV_STAT'] = 'ONTIME'
                extended_data.append(extended_row)
    df_extended = pd.DataFrame(extended_data)
    return df_extended.reset_index(drop=True)
```

```
[62]: create_data_2nd_model_data_ext(new_data_df).to_csv("SYR_ORIGIN_3HOP.csv",
    ↪index=False)
```

```
[63]: create_data_2nd_model_data_ext(new_data_df, 1).to_csv("SYR_ORIGIN_1HOP.csv",
    ↪index=False)
```

```
[64]: create_data_2nd_model_data_ext(new_data_df, 2).to_csv("SYR_ORIGIN_2HOP.csv",
    ↪index=False)
```

```
[ ]:
```